Index

activation function, 52
    hyperbolic tangent, 53
    ReLU, 53
    sigmoid, 53
    softmax, 53
    step function, 52
active learning, 102
adaptive step size, 60
    adadelta, 119
    adam, 120
    momentum, 118
    running average, 118
autoencoder, 103
    variational, 111

backpropagation through time, 126
bagging, 85
bandit problem
    contextual bandit problem, 102
    k-armed bandit, 102
basis functions
    polynomial basis, 42
    radial basis, 45
boosting, 78

classification, 29
    binary classification, 29
clustering, 103
    evaluation, 108
    into partitions, 104
    k-means algorithm, 105
convolution, 64
convolutional neural network, 64
    backpropagation, 68
cross validation, 22

data
    training data, 6
distribution, 11
    independent and identically distributed, 6
dynamic programming, 93

empirical probability, 83
ensemble models, 78
error
    test error, 11
    training error, 11
error back-propagation, 57
estimation, 6
estimation error, 21
expectation, 89
experience replay, 100
exploding (or vanishing) gradients, 60

features, 13
filter, 64
    channels, 65
    filter size, 66
finite state machine, 124
fitted Q-learning, 100

gating network, 131
generative networks, 111
gradient descent, 23
    applied to k-means, 106
    applied to logistic regression, 38
    applied to neural network training, 54
    applied to regression, 26
    applied to ridge regression, 27
    convergence, 24
    learning rate, 23
    stopping criteria, 24

hand-built features
    binary code, 47
    factored, 47
    numeric, 46
    one-hot, 47
    text encoding, 47
    thermometer, 46
hierarchical clustering, 108
hyperparameter
    hyperparameter tuning, 22
hyperplane, 15
hypothesis, 10
    linear regression, 15

k-means
    algorithm, 106

formulation, 105
in feature space, 108
initialization, 107
kernel methods, 46

language model, 126
latent representation, 109
learning
from data, 6
learning algorithm, 21
linear classifier, 30
hypothesis class, 30
linear logistic classifier, 33
linear time-invariant systems, 125
linearly separable, 31
logistic linear classifier
prediction threshold, 34
long short-term memory, 131
loss function, 10

Machine Learning, 6
Markov decision process, 87
max pooling, 66
model, 7
learned model, 10
model class, 7, 12
no model, 11
prediction rule, 11

nearest neighbor models, 78
negative log-likelihood, 35
loss function, 36
multi-class classification, 39
neural network, 50
batch normalization, 62
dropout, 61
feed-forward network, 51
fully connected, 51
initialization of weights, 58
layers, 51
loss and activation function choices, 54
output units, 51
regularization, 61
training, 58
weight decay, 61
neuron, 50
non-parametric methods, 78

optimal action-value function, 91
output function, 124

parameter, 11
policy, 89

finite horizon, 89
infinite horizon, 89
discount factor, 90
optimal policy, 91
reinforcement learning, 95
stationary policy, 93

Q-function, 91

real numbers, 13
recurrent neural network, 123
recurrent neural network, 125
gating, 131
regression, 13
linear regression, 15
random regression, 16
ridge regression, 20
regularizer, 14
reinforcement
actor-critic methods, 99
reinforcement learning
exploration vs. exploitation, 102
policy, 95
reinforcement learning, 87, 95, 123
Laplace correction, 101
RL algorithm, 96
reward function, 87, 101

separator, 30
linear separator, 31
sequence-to-sequence mapping, 126
sign function, 30
sliding window, 100
spatial locality, 63
state machine, 123
state transition diagram, 124
stochastic gradient descent, 28
convergence, 28
structural error, 21
supervised learning, 7

tensor, 65
total derivative, 127
transduction, 126
transition function, 124
transition model, 87
translation invariance, 63
tree models
information gain, 84
tree models, 78
building a tree, 82
classification, 83

cost complexity, 83
pruning, 83
regression tree, 81
splitting criteria, 84
entropy, 84
Gini index, 84
misclassification error, 84

value function
calculating value function, 91
value function, 91

weight sharing, 65